

Report on lab assignment 3

Convolutional Neural Network

Majd Jamal

May 5, 2021

1 Main objectives and scope of the assignment

My major goals in the assignment were

- to implement and train a convolutions neural network.
- use deep learning for solving natural language processing problems.
- to find a solution for unbalanced data set and use it for training.

I implemented a convolutional neural network from scratch using only NumPy. The architecture is based on two convolutional layers and one fully connected layer, with activation functions ReLU and Softmax. Moreover, I found a solution for training with unbalanced data.

2 Method

The network was built from scratch using Python 3 and packages: NumPy and Matplotlib. The code and plots are found in Appendix Code.

3 Result

(i) Gradients check

After implementing the required functions, I conducted three experiments to compare the analytic gradients with the numerical ones. The code for these experiments is shown in Appendix Code `analyzeGradient()`, on page 22. The absolute difference between numerical and analytic gradients was small, i.e., ($< 1e - 6$).

Few data points & small layer sizes

With 5 data points and parameters $\mathbf{n1} = 5$, $\mathbf{n2} = 5$, $\mathbf{k1} = 5$, $\mathbf{k2} = 3$, the difference was:

Weight (fully connected layer) Gradient: $9.88e-10$

Filter 2 (second convolutional layer) Gradient: $5.80e-11$

Filter 1 (first convolutional layer) Gradient: $7.33e-11$

Few data points & bigger layers

With **5 data points** and parameters $\mathbf{n1} = 20$, $\mathbf{n2} = 20$, $\mathbf{k1} = 5$, $\mathbf{k2} = 3$, the difference was:

Weight Gradient: $2.19e-10$

Filter 2 Gradient: $5.99e-11$

Filter 1 Gradient: $5.57e-11$

More data points & bigger layers

With **100 data points** and parameters $\mathbf{n1} = 20$, $\mathbf{n2} = 20$, $\mathbf{k1} = 5$, $\mathbf{k2} = 3$, the difference was:

Weight Gradient: $2.54e-09$

Filter 2 Gradient: $1.23e-09$

Filter 1 Gradient: $3.18e-09$

(ii) Compensating for unbalanced data

I accounted for unbalanced data by selecting six random points from each class. This method was similar to what the teacher used, as stated in the instruction paper.

(iii) Performance

The training data consists of 17204 data points. With a batch size of 108, it is required at least 125 epochs to compute ≥ 20000 update steps.

With unbalanced data

Settings for the training were:

n1 = 20, n2 = 20, k1 = 5, k2 = 3, eta = 0.001, roh = 0.9, epochs = 130, n_batches = 100

Score: 5.56%

The confusion matrix on the validation set is found in Appendix Figure 2. Moving Average Loss function on the validation set is found in Appendix Figure 3. Original Loss function on the validation set is found in Appendix Figure 4.

Comment: Due to learning with momentum, the loss curve oscillated, which leads to a disturbing pattern. I created a loss curve with 100 points moving average to increase interpretability.

With balanced data

Settings for the training were:

n1 = 20, n2 = 20, k1 = 5, k2 = 3, eta = 0.001, roh = 0.9, epochs = 130, n_batches = 108

Score: 41.2%

The confusion matrix on the validation set is found in Appendix Figure 5. Loss function on the validation set is found in Appendix Figure 6.

(iv) Best network

The intuition was to use the same setting as previous experiments but stops when the validation loss starts increasing, avoiding overfitting. This happens after about three epochs, as seen in figure 6. The settings was:

n1 = 20, n2 = 20, k1 = 5, k2 = 3, eta = 0.001, roh = 0.9, epochs = 3, n_batches = 108

, and the score was:

Score: 47.2%

The confusion matrix on the validation set is found in Appendix Figure 7. Loss function on the validation set are found in Appendix Figure 8.

(v) Compensating for unbalanced data

Q. State whether you implemented the efficiency gains in Background 5 .

I did not have time to improve my training speed, but I will improve it after the submission! However, I tried to pre-compute the first M_x -layer, but my computer stopped the computations because of memory issues. The computer used for this project has 8GB RAM, and the terminal broke after 3000/17402 computation, stating "zsh killed".

(vi) Testing

The names that were used for the prediction test was: [Linda, Per, Majd, Alba, Steve] , and the classifications were,

Linda: [German 27.8%, Polish 14.7%, Japanese 11.9%, Italian 10.9%, Czech 10.2%]

Per: [Chinese 58.7%, Korean 9.6%, French 7.5%, Czech 5.4%, Vietnamese 4.6%]

Majd: [Scottish 75.9%, Polish 6.9%, Czech 4.8%, Arabic 4.2%, German 2.5%]

Alba: [Spanish 23.7%, Italian 20.7%, Korean 18.7%, Irish 13.2%, Portuguese 4.47%]

Steve: [Italian 35.3%, Czech 15.9%, French 13.7%, Spanish 11.3%, German 8.8%]

4 Discussion

The performance was very poor when training with unbalanced data, and this was expected. We can clearly see by looking at the confusion matrix in Appendix Figure 2 that the network learned the dominant class, which is label 15. See figure 1 for a distribution of the labels.

The performance became much better when accounting for the unbalanced data. As seen in Figure 5, confusion matrix shows a colored diagonal, indicating that predicted labels matched true labels. The score also increased from 5% to 47.2%.

When testing the network with five names, it was found that classifications for Linda, Alba, and Steve were correct. For example, Linda could indeed be a german woman or polish. Alba is a Latin name used in Spanish and Italian-speaking countries.

This assignment showed that it is important to account for unbalanced data when training a deep network. For further developments and iterations, I will make my program much faster, because it took a very long time to train the network with 120-130 epochs, approximately six hours. An explanation for this slow training speed is that I used Kron multiplications when generating the MX-matrix.

5 Appendix

5.1 Result

(iii)




Figure 1: **Title:** Distribution of labels. **Comment:** We can see that label nr. 15 occurs more than the other classes.




Figure 2: **Title:** Confusion Matrix when training the network with unbalanced data. X-axis shows true labels and Y-axis shows predicted labels. Settings were: $n1 = 20$, $n2 = 20$, $k1 = 5$, $k2 = 3$, $\eta = 0.001$, $\rho = 0.9$, $\text{epochs} = 130$, $\text{n_batches} = 100$. **Comment:** We can clearly see that all the data points were classified as label 15, which is the dominant label in the distribution.




Figure 3: **Title:** 100 Moving Average Loss function on the validation set when training with unbalanced data. **Settings were:** $n1 = 20$, $n2 = 20$, $k1 = 5$, $k2 = 3$, $\eta = 0.001$, $\rho = 0.9$, $\text{epochs} = 130$, $\text{n_batches} = 100$. **Comment:** The loss function decreased steadily, indicating that the backpropagation algorithm is written correctly. However, the loss function eventually converged to approx 2.89 before 10 000 update steps, indicating that the network learned to classify points only as the dominant class.




Figure 4: **Title:** Loss function on the validation set when training with unbalanced data. **Settings were:** $n1 = 20$, $n2 = 20$, $k1 = 5$, $k2 = 3$, $\eta = 0.001$, $\rho = 0.9$, $\text{epochs} = 130$, $\text{n_batches} = 100$. **Comment:** The loss function oscillated heavily because of momentum learning. However, the curve had a downward moving pattern, indicating that the backpropagation was working properly.




Figure 5: **Title:** Confusion Matrix when training the network with balanced data. X-axis shows true labels and Y-axis shows predicted labels. Settings were: $n1 = 20$, $n2 = 20$, $k1 = 5$, $k2 = 3$, $\eta = 0.001$, $\rho = 0.9$, epochs = 130, n_batches = 108. **Comment:** We can see a diagonal on the confusion matrix, indicating that predictions were correct.

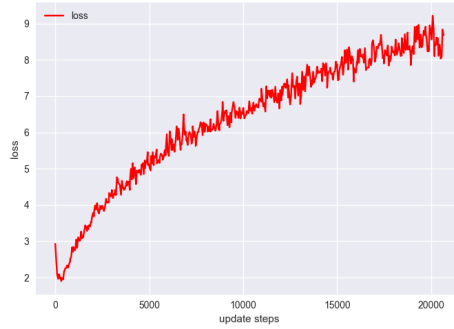


Figure 6: **Title:** 100 Loss function on the validation set when training with balanced data. **Settings were:** $n1 = 20$, $n2 = 20$, $k1 = 5$, $k2 = 3$, $\eta = 0.001$, $\rho = 0.9$, epochs = 130, n_batches = 108. **Comment:** The validation loss function decreased in the beginning, but started to increase after 200-300 update steps, because the network began overfitting.




Figure 7: **Title:** Confusion Matrix when training the best network. X-axis shows true labels and Y-axis shows predicted labels. Settings were: $n1 = 20$, $n2 = 20$, $k1 = 5$, $k2 = 3$, $\eta = 0.001$, $\rho = 0.9$, $\text{epochs} = 3$, $n_batches = 108$. **Comment:** We can clearly see a diagonal, meaning that many points were classified as their true class.

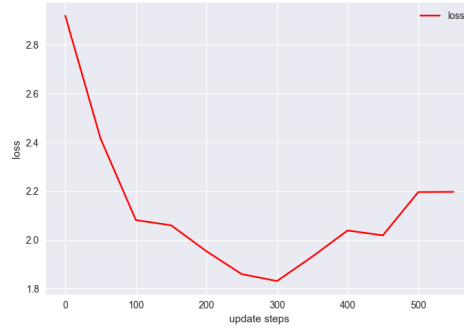


Figure 8: **Title:** Loss function on the validation set from training the best network. Settings were: $n1 = 20$, $n2 = 20$, $k1 = 5$, $k2 = 3$, $\eta = 0.001$, $\rho = 0.9$, $\text{epochs} = 3$, $n_batches = 108$.

5.2 Code

```
1  __author__ = "Majd Jamal"
2
3  import numpy as np
4  import matplotlib.pyplot as plt
5  import re
6
7
8  ##=====
9  ## Load data and seperate it into
10 ## names (string) and labels (int)
11 ##=====
12 file = 'dataset/ascii_names.txt'
13 fid = np.loadtxt(file, delimiter='\t', dtype=str)
14 Npts = fid.size
15
16 names = np.zeros(Npts).astype('str')
17 ys = np.zeros(Npts)
18
19 for i in range(Npts):
20     name, label = fid[i].split(' ')
21
22     names[i] = name.lower()
23     ys[i] = label
24
25
26 ##=====
27 ## Measure:
28 ##     number of unique unique characters
29 ##     length of the longest word
30 ##     number of classes
31 ##=====
32 doc = open(file, 'r').read().lower().replace("\n", "").replace(" ", "")
33
34 longest_word = max(names, key = len)    # Get longest string
35 characters = re.sub(r'\d+', '', doc) # Remove digits
36 unique_characters = ''.join(set(characters))    # Extract unique characters
37
38 NUnique = len(unique_characters) #Number of unique characters
39                                     #in the document
40 NLongest = len(longest_word)    #lenght of longest name
41 NClasses = np.max(ys).astype(int) #number of classes
42
43
44 ##=====
45 ## Create a lexicon of char2ind, i.e,
46 ## character to index-notation
47 ##=====
```

```

48 char2ind = {}
49
50 for i in range(NUnique):
51     char2ind[unique_characters[i]] = i
52
53
54 ##=====
55 ## Convert names to a matrix representation
56 ## i.e, the data point matrix X with shape = (Ndim, Npts)
57 ##=====
58 names2matrix = np.zeros((NUnique * NLongest, Npts)) #also known as, X
59
60 for i in range(Npts):
61     name = names[i]
62     name2vec = np.zeros((NUnique, NLongest))
63
64     for j in range(len(name)):
65
66         curr_char = name[j]
67         ind = char2ind[curr_char]
68
69         name2vec[ind][j] = 1
70
71     if i == 0:
72         xoriginal = name2vec
73         xoriginal_flatten = name2vec.flatten(order = 'F')
74         np.save('xoriginal.npy', xoriginal)
75         np.save('xoriginal_flatten.npy', xoriginal_flatten)
76
77     names2matrix[:, i] = name2vec.flatten(order = 'F')
78
79
80 ##=====
81 ## Creating one hot vector matrix
82 ## with shape = (Nclasses, Npts)
83 ##=====
84 Y = np.zeros((ys.size, ys.max().astype(int)+1))
85 Y[np.arange(ys.size), np.reshape(ys, (1,-1)).astype(int)] = 1
86 Y = Y.T
87 Y = np.delete(Y, 0, 0)
88
89
90 ##=====
91 ## Save files
92 ##=====
93 np.save('final/names.npy', names)
94 np.save('final/X.npy', names2matrix)
95 np.save('final/Y.npy', Y)
96 np.save('final/ys.npy', ys)
97 np.save('final/dims.npy', np.array([NUnique, NLongest, NClasses]))

```

```

98
99
100 ##-----
101 ## Utils
102 ##-----
103
104 class Data:
105     """Object to store data
106     """
107     def __init__(self,
108                 X_train, Y_train, y_train,
109                 X_val, Y_val, y_val,
110                 X, Y, y, names,
111                 NUnique, NLongest, NClasses):
112
113         self.X_train = X_train
114         self.Y_train = Y_train
115         self.y_train = y_train
116
117         self.X_val = X_val
118         self.Y_val = Y_val
119         self.y_val = y_val
120
121         self.X = X
122         self.Y = Y
123         self.y = y
124         self.names = names
125
126         self.NUnique = NUnique
127         self.NLongest = NLongest
128         self.NClasses = NClasses
129
130     def getData():
131         """Load and seperate data into
132         a training and validation set, and
133         return an object of the data.
134         """
135         X = np.load('data/final/X.npy')
136         Y = np.load('data/final/Y.npy')
137         y = np.load('data/final/ys.npy')
138
139         names = np.load('data/final/names.npy')
140
141         dims = np.load('data/final/dims.npy')
142
143         all_indices = np.arange(X.shape[1])
144         validation_indicies = np.loadtxt('data/dataset/Validation_Inds.txt').astype(int)
145         training_indicies = np.delete(all_indices, validation_indicies)
146
147

```

```

148     data = Data(
149         X_train = X[:, training_indicies] ,
150         Y_train = Y[:, training_indicies],
151         y_train = y[training_indicies],
152         X_val = X[:, validation_indicies],
153         Y_val = Y[:, validation_indicies],
154         y_val = y[validation_indicies],
155         X = X,
156         Y = Y,
157         y = y,
158         names = names,
159         NUnique = dims[0],
160         NLongest = dims[1],
161         NClasses = dims[2],
162     )
163
164     return data
165
166
167 class Params:
168     """ Object that are used to pass arguments to the deep network. """
169
170     def __init__(self, n1, n2, k1, k2, eta, roh, epochs, n_batches):
171
172         self.n1 = n1
173         self.n2 = n2
174
175         self.k1 = k1
176         self.k2 = k2
177
178         self.eta = eta
179         self.roh = roh
180
181         self.epochs = epochs
182         self.n_batches = n_batches
183
184     def softmax(x):
185         """ Standard definition of the softmax function """
186         e_x = np.exp(x - np.max(x))
187         return e_x / np.sum(e_x, axis=0)
188
189     def ReLU(x):
190         """ Standard definition of the ReLU function """
191         return np.maximum(x, 0)
192
193     def BatchCreator(j, n_batches):
194         """ Generates indices for the mini_batch,
195         given an iteration step and number of data points
196         in each mini_batch.
197         :param j: iteration step

```

```

198         :param n_batches: number of data points in a batch
199         :return ind: indicies for the mini batch
200         """
201         j_start = (j-1)*n_batches + 1
202         j_end = j*n_batches + 1
203         ind = np.arange(start= j_start, stop=j_end, step=1)
204         return ind
205
206     def vecX(x, d, nlen):
207         """ Vectorizes an x-data point in a similar fashion to Matlab.
208             :param x: data point, with shape = (Ndim, )
209             :param d: height of the original shape
210             :param nlen: width of the original shape
211             :return: vectorized version of data point x
212         """
213         return x.reshape((d, nlen)).flatten(order = 'F')
214
215     def vecF(F):
216         """ Vectorizes an a filter in a similar fashion to Matlab.
217             :param F: Filter, with shape (#filters, width, height)
218             :return: vectorized filter in 1D
219         """
220         nf,_,_ = F.shape
221         for filter in range(nf):
222
223             if filter == 0:
224                 F_flattened = F[filter].flatten(order = 'F')
225             else:
226                 F_flattened = np.hstack((F_flattened, F[filter].flatten(order = 'F')))
227
228         return F_flattened
229
230     def plotter(X, Y):
231         """ Plotting the loss function.
232             :param X: x-coordinates, which are used as update steps
233             :param Y: loss values
234         """
235         plt.style.use('seaborn')
236         plt.xlabel('update steps')
237         plt.ylabel('loss')
238         plt.plot(X, Y, color = 'red', label = 'loss')
239         plt.legend()
240         plt.savefig('result/loss')
241         plt.close()
242
243
244     ##=====
245     ## Model
246     ##=====
247

```

```

248
249 class ConvNet:
250     """ Convolutional Neural Network.
251         """
252     def __init__(self):
253
254         self.nlen = None    # Width of X
255         self.dx = None     # Height of X
256
257         self.F1 = None
258         self.F2 = None
259         self.W = None
260         self.MF1 = None
261         self.MF2 = None
262
263         self.counts = None # Count occurrences of each label,
264                             # e.g. {'0': 400 ...}
265
266         self.nlen1 = None  # Width of S1
267         self.preMX = []
268         self.losses = [[],[]] #Array to store validation losses
269         self.dL_dX = (0,0,0) #Memory used for momentum computations,
270                             #dW(t-1), dF2(t-1), dF1(t-1)
271
272     def MakeMFMatrix(self, F, nlen):
273         """ Creates the MF-matrix used for convolution operations.
274             :params F: A 3D filter with shape = (n_filters, height, width)
275             :params nlen: Width of the data points that
276                 goes through the filter layer.
277             :return MF: The MF matrix
278             """
279         nf,d,k = F.shape
280
281         zero = np.zeros((nf, nlen))
282
283         for filter in range(nf):
284             if filter == 0:
285                 VF = F[filter].flatten(order = 'F')
286             else:
287                 VF = np.vstack((VF, F[filter].flatten(order = 'F')))
288
289         MF = np.zeros(((nlen - k + 1)*nf, nlen*d))
290
291         step = 0
292         Nelements = VF[0].size
293
294         for i in range((nlen - k + 1)):
295
296             for j in range(nf):
297

```

```

298         ind = j + i*nf
299         MF[ind, step:Nelements + step] = VF[j]
300
301         step += d
302
303     return MF
304
305 def MakeMXMatrix(self, x_input, nf, d, k, dx, nlen):
306     """ Creates the MX-matrix used for convolution operations.
307     :params x_input: A 1D data point with shape = (height*width, )
308     :params nf: Number of filters in the layer
309     :params d: Height of the filter that are used in the convolutional layer
310     :params k: Width of the filter
311     :params dx: Original height of x_input
312     :params nlen: Original width of x_input
313     :return MX: The MX matrix
314     """
315     x_input = x_input.reshape((dx, nlen), order='F')
316
317     I_nf = np.identity(nf)
318
319     MX = np.zeros((
320         (nlen-k+1)*nf,
321         (k*nf*d)
322     ))
323
324     for i in range(nlen-k+1):
325
326         vec = x_input[:, i:i+k].T.flatten()
327
328         vec = np.kron(I_nf, vec)
329
330         if i == 0:
331             MX = vec
332
333         else:
334             MX = np.vstack((MX, vec))
335
336     return MX
337
338 def forward(self, X, MF1, MF2, W):
339     """ Computes the forward pass.
340     :param X: Data matrix, shape = (Ndim, Npts)
341     :param MF1: Filters for the first convolutional layer
342     :param MF2: filters for the first convolutional layer
343     :param W: Weights for the fully connected layer
344     :return X1: Values after the first ConvLayer
345     :return X2: Values after the second ConvLayer
346     :return P: Final predictions, shape = (Nout, Npts)
347     """

```



```

348     X1 = ReLU(MF1 @ X)
349     X2 = ReLU(MF2 @ X1)
350     S = W@X2
351     P = softmax(S)
352     return X1, X2, P
353
354 def backward(self, S1, S2, P, W, XBatch, YBatch):#, ind):
355     """ Computes the backward pass.
356     :param S1: Values after the first ConvLayer
357     :param S2: Values after the second ConvLayer
358     :param P: Final predictions of the network
359     :param W: Weights for the fully connected layer
360     :param XBatch: X mini batch
361     :param YBatch: Y mini batch
362     :return dW: Gradients for the fully connecte layer
363     :return dF2: Gradients for the second convolutional layer filters
364     :return dF1: Gradients for the first convolutional layer filters
365     """
366     _, Npts = P.shape
367
368     G = - (YBatch - P)
369
370     dW = 0
371
372     for j in range(Npts):
373         g = G[:, j].reshape(-1, 1)
374         s2 = S2[:, j].reshape(-1, 1)
375         y = YBatch[:, j].argmax()
376         #py = (1/self.counts[y]) * (1/18)
377
378         dW += g@s2.T #* py
379
380     dW /= Npts
381
382     G = W.T @ G
383     G = G * np.where(S2 > 0, 1, 0)
384
385     nf, _, _ = self.F2.shape
386     dF2 = 0
387
388     for j in range(Npts):
389         g = G[:, j].reshape(-1, 1)
390         x = S1[:, j].reshape(-1, 1)
391         mx = self.MakeMXMatrix(x, *self.F2.shape, nf, self.nlen1)
392         v = g.T @ mx
393         y = YBatch[:, j].argmax()
394         #py = (1/self.counts[y]) * (1/18)
395
396         dF2 += v #* py

```

```

397
398     dF2 /= Npts
399
400     G = self.MF2.T @ G
401     G = G * np.where(S1 > 0, 1, 0)
402
403
404     nf,d,k = self.F1.shape
405
406     dF1 = 0
407     for j in range(Npts):
408         g = G[:, j].reshape(-1, 1)
409         x = XBatch[:, j].reshape(-1, 1)
410         mx = self.MakeMXMatrix(x, *self.F1.shape, self.dx, self.nlen)
411         v = g.T @ mx
412         y = YBatch[:, j].argmax()
413         #py = (1/self.counts[y]) * (1/18)
414
415         dF1 += v * py
416
417     dF1 /= Npts
418
419     return dW, dF2, dF1
420
421 def update(self, dW, dF2, dF1, eta, rho):
422     """ Updates the weights and filters.
423     :param dW: Gradients for the fully connecte layer
424     :param dF2: Gradients for the second convolutional layer filters
425     :param dF1: Gradients for the first convolutional layer filters
426     :param eta: Learning rate
427     :param rho: Momentum constant
428     """
429
430     nf, d, k = self.F2.shape
431     nf1, d1, k1 = self.F1.shape
432
433     # -= gradient * learning rate + previous_gradient * momentum
434     self.W -= (dW * eta
435               + self.dL_dX[0] * rho)
436
437     self.F2 -= (dF2.reshape((d,k, nf), order='F').transpose([2,0,1]) * eta
438               + self.dL_dX[1].reshape((d,k, nf), order='F').transpose([2,0,1]) * rho)
439
440     self.F1 -= (dF1.reshape((d1,k1, nf1), order='F').transpose([2,0,1]) * eta
441               + self.dL_dX[2].reshape((d1,k1, nf1), order='F').transpose([2,0,1]) * rho)
442
443 def ComputeCost(self, X, Y, MF1, MF2, W, F1, F2):
444     """ Computes the loss function.
445     :param X: Data matrix, shape = (Ndim, Npts)
446     :param Y: One hot encoded labels, shape =(Nout, Npts)

```

```

447         :param MF1: Filter matrix for the first convolutional layer
448         :param MF2: Filters for the second convolutional layer
449         :param W: Weight for the fully connected layer.
450         """
451         _, P = self.forward(X, MF1, MF2, W)
452         _, Npts = P.shape
453
454         loss = 0
455
456         for j in range(Npts):
457
458             y = Y[:, j]
459             p = P[:, j]
460             ind = y.argmax()
461             #py = (1/self.counts[ind]) * (1/18)
462             loss -= np.log(y.T @ p) #* py
463
464         loss /= Npts
465
466         return loss
467
468     def ComputeAccuracy(self, P, y):
469         """ Computes a score indicating the network accuracy.
470         :param P: Probabilities, shape = (Nout, Npts)
471         :param y: labels, shape = (Npts, )
472         :return: error rate. low is better.
473         """
474         out = np.argmax(P, axis=0).reshape(1,-1)
475         #np.savetxt('out.txt', out.astype(int))
476         return 1 - np.mean(np.where(y==out, 0, 1))
477
478     def ComputeGradsNumSlow(self, X, Y, W, F2, F1, h):
479         """ Computes numerical gradients.
480         """
481         MF2 = self.MakeMFMatrix(F2, self.nlen1)
482         MF1 = self.MakeMFMatrix(F1, self.nlen)
483
484         grad_W1 = np.zeros(W.shape)
485         for i in range(W.shape[0]):
486             for j in range(W.shape[1]):
487                 W1_try = np.array(W)
488                 W1_try[i,j] -= h
489                 c1 = self.ComputeCost(X, Y, MF1, MF2, W1_try, F1, F2)
490
491                 W1_try = np.array(W)
492                 W1_try[i,j] += h
493                 c2 = self.ComputeCost(X, Y, MF1, MF2, W1_try, F1, F2)
494
495                 grad_W1[i,j] = (c2 - c1) / (2 * h)
496

```

```

497     grad_F2 = np.zeros(F2.shape)
498     for k in range(F2.shape[0]):
499         for i in range(F2.shape[1]):
500             for j in range(F2.shape[2]):
501
502                 F2_try = np.array(F2)
503                 F2_try[k, i, j] -= h
504                 MF2_try = self.MakeMFMatrix(F2_try, self.nlen1)
505                 c1 = self.ComputeCost(X, Y, MF1, MF2_try, W, F1, F2)
506
507                 F2_try = np.array(F2)
508                 F2_try[k, i, j] += h
509                 MF2_try = self.MakeMFMatrix(F2_try, self.nlen1)
510                 c2 = self.ComputeCost(X, Y, MF1, MF2_try, W, F1, F2)
511
512                 grad_F2[k, i, j] = (c2 - c1) / (2 * h)
513
514     grad_F1 = np.zeros(F1.shape)
515     for k in range(F1.shape[0]):
516         for i in range(F1.shape[1]):
517             for j in range(F1.shape[2]):
518
519                 F1_try = np.array(F1)
520                 F1_try[k, i, j] -= h
521                 MF1_try = self.MakeMFMatrix(F1_try, self.nlen)
522                 c1 = self.ComputeCost(X, Y, MF1_try, MF2, W, F1, F2)
523
524                 F1_try = np.array(F1)
525                 F1_try[k, i, j] += h
526                 MF1_try = self.MakeMFMatrix(F1_try, self.nlen)
527                 c2 = self.ComputeCost(X, Y, MF1_try, MF2, W, F1, F2)
528
529                 grad_F1[k, i, j] = (c2 - c1) / (2 * h)
530
531     return grad_W1, grad_F2, grad_F1
532
533 def TestMFandMX(self):
534     """ This function test if the implementation of
535     MX and MF is correct.
536     """
537     X_test = np.arange(1*6*4) + 1
538     X_test = X_test.reshape(1,6,4)
539     #print(X_test)
540     X_input = X_test.flatten(order = 'F')
541     #print(X_input)
542
543     F_test = np.arange(4*6*3) + 1
544     F_test = F_test.reshape(4,6,3)
545     #print(F_test)
546

```

```

547     #nf,d,k = F_test.shape
548
549     MF_test = self.MakeMFMatrix(F_test, 4)
550     #print(MF_test)
551     print(X_input)
552     MX_test = self.MakeMXMatrix(X_input, *F_test.shape, 6, 4)
553     print(MX_test)
554
555     s1 = MF_test @X_input
556     s2 = MX_test @ vecF(F_test)
557
558     print(np.all(s1 == s2)) # >>> True
559
560 def debug(self):
561     """ This section test if the network functions
562     can reproduce vectors from DebugInfo.mat.
563     """
564
565     import scipy.io
566     d = scipy.io.loadmat('utils/DebugInfo.mat')
567
568     debug_x = d['x_input']
569     debug_F = d['F']
570     debug_vecF = d['vecF']
571     debug_vecS = d['vecS']
572     debug_S = d['S']
573     debug_X = d['X_input']
574
575     dx, nlen = debug_X.shape
576     #print(debug_F.shape)
577     debug_F = debug_F.transpose([2,0,1])
578     nf, d, k = debug_F.shape
579
580     MF = self.MakeMFMatrix(debug_F, nlen)
581     S1 = MF @ debug_x
582
583     print(S1[:, 0] == debug_vecS[:, 0])
584
585     MX = self.MakeMXMatrix(debug_x, d, k, nf, dx, nlen)
586     S2 = MX @debug_vecF
587     print(S2[:, 0] == debug_vecS[:, 0])
588
589     my_S = S2.reshape(debug_S.shape)
590
591     print(my_S == debug_S)
592
593 def plotLabelDist(self):
594     """ Plots label distribution.
595     """
596     import matplotlib.pyplot as plt

```

```

597     plt.style.use('seaborn')
598     plt.bar([str(i) for i in range(1,19)], self.counts, color='#ff7f4f')
599     plt.xlabel('label')
600     plt.ylabel('occurences')
601     plt.show()
602
603     def AnalyzeGradients(self, X, Y):
604         """ Computes and prints the difference between
605         numerical and analytical gradients.
606         :param X: Data matrix, shape = (Ndim, Npts)
607         :param Y: One hot matrix, shape = (Nout, Npts)
608         """
609         XBatch = X[:, :100]
610         YBatch = Y[:, :100]
611
612         S1, S2, P = self.forward(XBatch, self.MF1, self.MF2, self.W)
613
614         grad_an, grad_an_F2, grad_an_F1 = self.backward(S1, S2, P, self.W, XBatch, YBatch)
615
616         grad_num, grad_num_F2, grad_num_F1 = self.ComputeGradsNumSlow(XBatch, YBatch, self
617
618         nf, d, k = self.F2.shape
619         nf1, d1, k1 = self.F1.shape
620         grad_an_F1 = grad_an_F1.reshape((d1,k1, nf1), order='F').transpose([2,0,1])
621         grad_an_F2 = grad_an_F2.reshape((d,k, nf), order='F').transpose([2,0,1])
622
623         diff = np.sum(np.abs(np.subtract(grad_an, grad_num)))
624         diff /= np.sum(np.add(np.abs(grad_an), np.abs(grad_num)))
625         print('W: ', diff)
626
627         diff = np.sum(np.abs(grad_an_F2 - grad_num_F2))
628         diff /= np.sum(np.abs(grad_an_F2) + np.abs(grad_num_F2))
629         print('F2: ',diff)
630
631         diff = np.sum(np.abs(np.subtract(grad_an_F1, grad_num_F1)))
632         diff /= np.sum(np.add(np.abs(grad_an_F1), np.abs(grad_num_F1)))
633         print('F1: ',diff)
634
635     def getLoss(self):
636         """ Returns the loss function.
637         """
638         return self.losses
639
640     def getWeights(self):
641         """ Returns weights
642         """
643         return self.F1, self.F2, self.W
644
645     def MakeConfusionMatrix(self, P, true, Nout):
646

```

```

647     """ Creates a confusion matrix of predictions and saves it in result/
648     :param P: Final predicted probabilities
649     :param true: true labels
650     :param Nout: Number of unique classes
651     """
652     pred = np.argmax(P, axis=0)
653     CM = np.zeros((Nout, Nout))
654     _, counts = np.unique(true, return_counts = True)
655     for i in range(true.size):
656         true_y = true[i].astype(int)
657         pred_y = pred[i].astype(int)
658         CM[true_y, pred_y] += 1 / counts[true_y]
659
660     import matplotlib.pyplot as plt
661     plt.title('Confusion Matrix')
662     im = plt.imshow(CM, cmap = 'Blues')
663     plt.xlabel('Label')
664     plt.ylabel('Label')
665     plt.yticks(np.arange(18), np.arange(1, 19))
666     plt.xticks(np.arange(18), np.arange(1, 19))
667     bar = plt.colorbar(im)
668     #plt.show()
669     plt.savefig('result/CM')
670     plt.close()
671
672     def name2vec(self, name):
673         """ Create a flattened vector representation of names.
674         :param name: string
675         :return: Vector representation with shape = (Height*Width, )
676         """
677
678         name2vec = np.zeros((self.dx, self.nlen))
679
680         for j in range(len(name)):
681
682             curr_char = name[j]
683             ind = self.char2ind.item().get(curr_char)
684
685             name2vec[ind][j] = 1
686
687         return name2vec.flatten(order = 'F')
688
689     def fit(self, data, p):
690         """ This function is called to start training.
691         :param data: Object containing training and validation data
692         :param p: Object containing parameters used for training, i.e.
693         epochs, n_batch, eta, etc.
694         """
695
696     ##

```

```

697     ## Data
698     ##
699     X_train = data.X_train
700     Y_train = data.Y_train
701     y_train = data.y_train - 1
702
703     X_val = data.X_val
704     Y_val = data.Y_val
705     y_val = data.y_val - 1
706
707     self.char2ind = np.load('data/final/char2ind.npy', allow_pickle=True)
708
709     ##
710     ## Parameters
711     ##
712     Ndim, Npts = X_train.shape
713     self.dx, self.nlen = data.NUnique, data.NLongest
714     Nout, _ = Y_train.shape
715     epochs = p.epochs
716     n_batches = p.n_batches
717
718     d = data.NUnique           #height of F1
719     nlen = data.NLongest       #width of X
720     nlen1 = nlen - p.k1 + 1    #width of X1
721     nlen2 = nlen1 - p.k2 + 1   #width of X2
722
723     self.nlen = nlen
724     self.nlen1 = nlen1
725
726     ##
727     ## Filters & Weights using He-initialization
728     ## source: https://towardsdatascience.com/
729     ## weight-initialization-techniques-in-neural-networks-26c649eb3b78
730     ## F1, F2, W, MF1, and MF2
731     ##
732     self.F1 = np.random.randn(p.n1, d, p.k1) * np.sqrt(2/d)
733     self.F2 = np.random.randn(p.n2, p.n1, p.k2) * np.sqrt(2/(d*p.k1))
734     self.W = np.random.randn(Nout, (p.n2 * nlen2)) * np.sqrt(2/(p.n1*p.k2))
735
736     self.MF1 = self.MakeMFMatrix(self.F1, nlen)
737     self.MF2 = self.MakeMFMatrix(self.F2, nlen1)
738     _, self.counts = np.unique(y_train, return_counts = True)
739     self.dL_dX = (np.zeros(self.W.shape),
740                  np.zeros(self.F2.size),
741                  np.zeros(self.F1.size))
742
743     ##
744     ## Debug
745     ##
746     #self.TestMFandMX() # Test implementation of MF and MX

```



```

747     #self.debug() # Take the Debug test
748     #self.AnalyzeGradients(X_train, Y_train) # Analyze gradients.
749
750     print('=== Settings === \n epochs: ', epochs, ' steps/epoch: ', round(Npts/n_batches))
751     print('=== Starting Training ===')
752
753     indices = [np.where(y_train == cl)[0] for cl in range(Nout)]
754
755     for i in range(epochs):
756         for j in range(round(Npts/n_batches)):
757             XBatch = 0
758             YBatch = 0
759
760             ##
761             ## Generate mini batches with equal
762             ## label distribution, i.e. accounting
763             ## for the unbalanced data set.
764             ##
765             for cls in range(Nout):
766                 rnd = np.random.randint(low=0, high=indices[cls].size, size=6)
767                 if cls == 0:
768                     XBatch = X_train[:, indices[cls][rnd]]
769                     YBatch = Y_train[:, indices[cls][rnd]]
770                 else:
771
772                     XBatch = np.hstack((XBatch, X_train[:, indices[cls][rnd]]))
773                     YBatch = np.hstack((YBatch, Y_train[:, indices[cls][rnd]]))
774
775
776             self.MF1 = self.MakeMFMatrix(self.F1, nlen)
777             self.MF2 = self.MakeMFMatrix(self.F2, nlen1)
778
779             S1, S2, P = self.forward(XBatch, self.MF1, self.MF2, self.W)
780
781             gradients = self.backward(S1, S2, P, self.W, XBatch, YBatch)#, ind)
782
783             self.update(*gradients, p.eta, p.roh)
784             self.dL_dX = gradients
785
786             # Evaluate the model after 50th update step
787             if j % 50 == 0:
788                 loss = self.ComputeCost(X_val, Y_val, self.MF1, self.MF2, self.W, self
789                 update_step = i * round(Npts/n_batches) + j
790                 print('loss: ', loss)
791                 self.losses[0].append(update_step)
792                 self.losses[1].append(loss)
793
794             print('Epoch: ', i + 1)
795             print('\n')
796

```

```

797         print('=== Training Completed ===')
798
799         ##
800         ## Network evaluation
801         ##
802         S1, S2, P = self.forward(X_val, self.MF1, self.MF2, self.W)
803         self.MakeConfusionMatrix(P, y_val, Nout)
804         acc = self.ComputeAccuracy(P, y_val)
805         print('Accuracy: ', acc)
806         #"""
807
808     def predict(self, name):
809         """ Predicts top 5 labels and their probabilities
810             for a data point.
811             :param name: string
812             :return out: top 5 predicted labels
813             :return prob: probabilities for the labels
814         """
815         X = self.name2vec(name)
816         _,_, P = self.forward(X, self.MF1, self.MF2, self.W)
817         out = P.argsort(axis= 0)[-5:][::-1]
818         prob = P[out]
819
820         return out, prob
821
822
823     ##=====
824     ## Experiments
825     ##=====
826
827
828     #=====
829     # Setup
830     #=====
831     data = getData()
832     params = Params(
833         n1 = 20, n2 = 20,
834         k1 = 5, k2 = 3,
835         eta = 0.001, roh = 0.9,
836         epochs = 3, n_batches = 108)
837
838     cnn = ConvNet()
839
840     #=====
841     # Training
842     #=====
843     cnn.fit(data, params)
844     weights = cnn.getWeights() #[F1, F2, W]
845     loss_ind, loss = cnn.getLoss()
846     plotter(loss_ind, loss)

```

```

847
848 #Save weights
849 np.save('data/weights/weights.npy', weights)
850 np.save('loss.npy', [loss_ind, loss])
851
852 #=====
853 # Prediction Test
854 #=====
855 print('===== Test =====')
856 test = ['linda', 'per', 'majd', 'alba', 'steve']
857 for name in test:
858     labels, probabilities = cnn.predict(name)
859     print('Name: ', name)
860     print('lbl ', labels)
861     print('pr ', probabilities)
862     print('\n')

```